

OPEN SOURCE

MIT-0

Agentic Coding Harness & Benchmarks

Cut agentic-coding **token spend** by choosing the harness and model from a measured **cost/quality frontier** -- on Amazon Bedrock or self-hosted.

CLAUDE CODE, PI, OMP & KIRO-CLI


AMAZON BEDROCK

SELF-HOSTED VLLM

LLM-AS-JUDGE

Executive briefing

swe-router in 70 seconds



The right model, handed to the developer for each task. The platform team benchmarks once, and spends less.

What you'll get out of this

A buying guide

A measured cost/quality frontier -- pick the model on evidence, not reputation.

Measured on your work

17 models over 21 real tasks from real repositories, each scored by an independent judge.

Your hosting, your call

One measurement across Amazon Bedrock and self-hosted GPUs -- decide on a level footing.

And the skill that reads the frontier for you: **swe-router** names the cheapest model that clears the bar for the task in front of you.

Agentic coding burns tokens **at scale**

Long-horizon, not one-shot

A task is tens to hundreds of LLM turns (~50-250, up to 300+), running 10 minutes to over an hour -- not a single prompt/response.

Prefill-heavy shape

Each turn replays a growing transcript as fresh input and emits a tiny edit. Real input:output runs **~150:1 to ~660:1** -- not the ~3:1 generic pricing assumes.

Cost hides in the choices

How many turns, how much context re-read, which harness, which model -- these swing the bill several-fold for similar quality.

Organizations are already at -- or heading toward -- **trillions of tokens per month, and coding agents drive a large and growing share of it. That is the bill this repo helps control.**

Public leaderboards can be saturated, and generic token pricing does not reflect this workload -- you need numbers on work that looks like yours.

Why it's different

Coding spend **doesn't look like app traffic**

App traffic

single-shot · small / volume models



Agentic coding

~50–250 turns · 10 min–1 hr+ · frontier models



■ input ■ output Same bar width - watch the output slice all but vanish for coding. That's where the cost is.

Generic pricing and leaderboards don't tell you the bill. Only benchmarking on real tasks does.

ORG A · FRONTIER-ONLY

Everything runs on a premium model



~90% of volume on premium. Nothing routes down. The bill scales with headcount × frontier price.

ORG B · GOVERNED MIX

Premium where it counts, open-weight for the rest




~70% of volume on open-weight models at a fraction of the per-token price - premium spread across two vendors and reserved for work that needs it.

■ Premium / frontier ■ Small proprietary ■ Open-weight, self-hosted


Same developers, same tasks - **very different bills**. The gap isn't discipline; it's whether anyone measured which tasks actually need a frontier model

Measure once, spend less on every task



The platform team owns the left half and runs it unattended. The developer never leaves their editor: the skill is already installed, engages on its own, and prints a recommendation before any work starts.

The cost/quality frontier (omp, /swe3)



17 models, 21 real tasks on mcp-gateway-registry-v2. Anthropic points are metered Bedrock bills, self-hosted points hardware-derived, so compare within a hosting basis.

Pick any model below the dashed line. There is a model *on* the line that scores higher **and** costs less. So there is never a reason to choose the one below.

swe-router, **in use**

[placeholder]

Screen recording: a developer starts a task, the skill engages on its own, prints its recommendation, and the developer switches model.

Replace with the terminal capture (animated GIF or MP4) before presenting.

Make it easy to do the right thing

THE FRONTIER

Continuous benchmarking - like evals - is desirable, not a burden: new models launch every week, so the frontier keeps moving.



The frontier only exists when all three come together - which is why **no single model provider can hand it to you**. You measure it on your own tasks.

NOW · RUN IT AS A HABIT

A framework + open-source implementation

The framework *and* the code to benchmark model × harness × hosting on your own real tasks. Re-run it regularly; hand your developers the result as an approved model/harness guideline.

NEXT · SHIFT LEFT SHIPS TODAY

An "swe-router" skill, inside any harness

Drops the frontier into your agentic-coding harness - whichever one you use - so the right model and harness are chosen automatically. Your developers stop carrying that decision.

Where the decision lives

Routing belongs **where the context lives**

APP TRAFFIC · GATEWAY-SIDE



AGENTIC CODING · SHIFTED LEFT INTO THE IDE



Same word, two different jobs. At the gateway it is **routing** - a classification made on the prompt alone, which is right for application traffic. For coding the context *is* the repository, so the choice is **decisioning**: a judgement the harness can only form after it has looked around and asked. In effect the decision **shifts left - out of the gateway and into the IDE**, where the agentic coding harness already runs.

Follows from the input-heavy profile earlier: ~50-250 turns replaying a growing transcript, not a single self-contained request. Gateways still own the things that *are* stateless - rate limits, spend caps, policy.

Take the model choice off the developer

Task + repo → tier → frontier lookup → model · harness · effort. A swe-router skill triages the task, takes the **cheapest model on the frontier that clears the bar**, and escalates a tier if the result falls short.

OPTION 1 · START HERE · ADVISE

Tell the developer, loudly

"For this task, switch to **this model at this reasoning effort**." The developer keeps their harness and can override it. No trust barrier - and every accepted call is evidence the routing works.

OPTION 2 · THEN · EXECUTE NEXT

Run it headless, hand back the result

The skill launches a **headless instance** of the chosen harness, at the chosen model and thinking effort, and completes the task. No manual switch, no decision left on the developer.

Sequence it: option 1, then option 2. Advisory mode earns the trust; move to headless once developers stop second-guessing the pick. Same frontier underneath - only the amount of automation changes.

Run it on your own code

Nothing here is specific to the example repository. Point the harness at your repos, run your model list through it, and the same generators build the frontier for **your** code. Then install **swe-router** so your developers spend that measurement without looking anything up.

1. Point at a repo

Any GitHub URL + a task description.
The example repo is the example, not
the contract.

2. Pick a path

Anthropic on Bedrock, open-weight on
Bedrock (LiteLLM), or self-hosted
vLLM.

3. Read your frontier

Score with the judge, plot cost against
quality, then ship the result as the skill.

[GitHub Repository](#)

[Install swe-router](#)

Scan to explore

The full harness, benchmarks, results, and methodology -- open source under MIT-0.

github.com/aarora79/agentic-coding-harness-benchmarks



Scan for the repository

HELD BACK

Backup

How to read the frontier, and the model-by-model detail behind it.

What is a **Pareto frontier**?

The set of models where no other model is both higher-scoring AND cheaper. Those are the only models worth considering. Everything else is *dominated* -- some frontier model beats it on both axes, so there is never a reason to pick it.

On the frontier

Non-dominated: to do better you must pay more, and to pay less you must accept a lower score. The real trade-off curve.

Dominated (skip it)

Example from our data: **kimi-k2.7-code (69.98, \$3.13/task)** is beaten by **qwen3.8-27b (78.48, \$1.47/task)** -- higher score *and* lower cost (both self-hosted). kimi-k2.7-code is off the frontier.

How to read it

Pick the quality your task needs (y-axis), then take the **leftmost model on the frontier** at that level -- the cheapest way to buy that quality.

Cost bases differ (metered Bedrock vs hardware-derived self-hosted), so we also draw the frontier within each hosting basis.

What the frontier tells a buyer

82.8

claude-opus-5

Top quality, \$11.95/task -- the premium tier for business-critical work.

80.1

glm-5.3

Best open-weight model at \$7.04/task; the self-hosted quality anchor.

78.5

qwen3.8-27b

Within 4.4 points of the top model at \$1.47/task -- an eighth of its cost.

A 27B open-weight model sits **4.4 points behind the best model at an eighth of the cost**. Which model you pick moves the bill more than anything else you control.

Pick by tier (omp, /swe3, 21 tasks)

Premium**claude-opus-5**

Business-critical / high-stakes changes. Highest quality (82.83) at \$11.95/task, finishing 21/21; wrong designs are expensive here.

Premium open-weight**glm-5.3**

Best self-hosted quality (80.11, \$7.04/task), 2.7 points off the top model. Standardize on it if you self-host your frontier tier.

Reliable workhorse**qwen3.8-27b**

Bulk of day-to-day self-hosted coding: 78.48 at \$1.47/task. It finished 20 of 21, so check the one it drops.

Most cost-effective**qwen3.6-35b**

Boilerplate, small fixes, scaffolding you'll review anyway: 59.24 at \$0.26/task, finishing 21/21.

A cheap model that sits off the frontier, or that does not finish the task, is no bargain.